


Math Review-1

For ENGN8535

Math for ENGN8535



Subjects

- Linear Algebra (Matrix computation, Matrix calculus)
- Probability theory (and statistical inference)
- Numerical Optimisation
- Python programming (Numpy)
- → Example applications in Data Analysis:
Ranking and Recommendation algorithms.

Some Calculus

Derivatives

- For a single valued function $f(x_1, \dots, x_n): \mathbb{R}^n \rightarrow \mathbb{R}$

- Partial derivative of x_i :
$$\frac{\partial f}{\partial x_i} = \lim_{h \rightarrow 0} \frac{f(x + he_i) - f(x)}{h}$$

- Gradient of f is
$$\nabla f(x) = \begin{pmatrix} \partial f / \partial x_1 \\ \partial f / \partial x_2 \\ \vdots \\ \partial f / \partial x_n \end{pmatrix}$$

- Directional derivatives: for $\|p\| = 1$

$$D(f(x), p) = \lim_{h \rightarrow 0} \frac{f(x + hp) - f(x)}{h} = \nabla f(x)^T p.$$

Hessian

- : the second order derivative of f

$$\nabla^2 f(x) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1 \partial x_1}(x) & \frac{\partial^2 f}{\partial x_1 \partial x_2}(x) & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n}(x) \\ \frac{\partial^2 f}{\partial x_2 \partial x_1}(x) & \frac{\partial^2 f}{\partial x_2 \partial x_2}(x) & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n}(x) \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1}(x) & \frac{\partial^2 f}{\partial x_n \partial x_2}(x) & \cdots & \frac{\partial^2 f}{\partial x_n \partial x_n}(x) \end{pmatrix}$$

- If f is twice continuously differentiable, Hessian is symmetric.

Taylor's theorem

- Mean value theorem: for $\alpha \in (0,1)$

$$f(x + p) = f(x) + \nabla f(x + \alpha p)^T p$$

- Taylor's theorem: for $\alpha \in (0,1)$

$$f(x + p) = f(x) + \nabla f(x)^T p + \frac{1}{2} p^T \nabla^2 f(x + \alpha p) p$$

Vector valued function

- For a vector valued function $r: \mathbb{R}^n \rightarrow \mathbb{R}^m$

$$r(x) = \begin{pmatrix} f_1(x_1, x_2, \dots, x_n) \\ f_2(x_1, x_2, \dots, x_n) \\ \vdots \\ f_m(x_1, x_2, \dots, x_n) \end{pmatrix}$$

- The Jacobian of r at x is



$$J(x) = \begin{pmatrix} \frac{\partial f_1}{\partial x_1}(x) & \frac{\partial f_1}{\partial x_2}(x) & \dots & \frac{\partial f_1}{\partial x_n}(x) \\ \frac{\partial f_2}{\partial x_1}(x) & \frac{\partial f_2}{\partial x_2}(x) & \dots & \frac{\partial f_2}{\partial x_n}(x) \\ \vdots & \vdots & & \vdots \\ \frac{\partial f_m}{\partial x_1}(x) & \frac{\partial f_m}{\partial x_2}(x) & \dots & \frac{\partial f_m}{\partial x_n}(x) \end{pmatrix} = \begin{pmatrix} \nabla f_1(x)^T \\ \nabla f_2(x)^T \\ \vdots \\ \nabla f_m(x)^T \end{pmatrix}$$

Convex function

- A function f is convex if for $\alpha \in [0,1]$

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$

- A function f is concave if $-f$ is convex



Lipschitz continuity

- A function f is said to be Lipschitz continuous on some set \mathcal{N} if there is a constant $L > 0$ such that

$$||f(x) - f(y)|| \leq L ||x - y|| \text{ for all } x, y \in \mathcal{N}$$

- If function f and g are Lipschitz continuous on \mathcal{N} , $f + g$ and fg are Lipschitz continuous on \mathcal{N} .

- ENGN8535 2021 S1: Math Review-1

Linear Algebra

Basic concepts

- **Vector** in \mathbb{R}^n is an ordered set of n real numbers.

- e.g. $v = (1, 6, 3, 4)$ is in \mathbb{R}^4

- A column vector:

$$\begin{pmatrix} 1 \\ 6 \\ 3 \\ 4 \end{pmatrix}$$

- A row vector:

- m -by- n **matrix** is an object in $\mathbb{R}^{m \times n}$ with m rows and n columns, each entry filled with a (typically) real number:

$$(1 \quad 6 \quad 3 \quad 4)$$

$$\begin{pmatrix} 1 & 2 & 8 \\ 4 & 78 & 6 \\ 9 & 3 & 2 \end{pmatrix}$$

Basic concepts

Vector norms: A norm of a vector $\|x\|$ is informally a measure of the “length” of the vector.

$$\|x\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$

- Common norms: L_1 , L_2 (Euclidean)

$$\|x\|_1 = \sum_{i=1}^n |x_i| \quad \|x\|_2 = \sqrt{\sum_{i=1}^n x_i^2}$$

- L_{infinity}

$$\|x\|_{\infty} = \max_i |x_i|$$

Vector

- A column vector $x \in \mathbb{R}^n$ is denoted as $x = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}$
- The transpose of x is $x^T = (x_1 \ x_2 \ \cdots \ x_n)$
- The inner product of $x, y \in \mathbb{R}^n$ is $x^T y = \sum_{i=1}^n x_i y_i$
- Vector norm

- 1-norm

$$\|x\|_1 = \sum_{i=1}^n |x_i|$$

- 2-norm

$$\|x\|_2 = \sqrt{x^T x} = \sqrt{\sum_{i=1}^n x_i^2}$$

- ∞ -norm

$$\|x\|_\infty = \max_{i=1 \dots n} |x_i|$$

Useful inequalities involving norms

Proposition 4.2. *The following inequalities hold for all $x \in \mathbb{R}^n$ (or $x \in \mathbb{C}^n$):*

$$\|x\|_{\infty} \leq \|x\|_1 \leq n\|x\|_{\infty},$$

$$\|x\|_{\infty} \leq \|x\|_2 \leq \sqrt{n}\|x\|_{\infty},$$

$$\|x\|_2 \leq \|x\|_1 \leq \sqrt{n}\|x\|_2.$$

Matrix

- A matrix $A \in \mathbb{R}^{m \times n}$ is

$$A = \begin{pmatrix} A_{11} & A_{12} & \cdots & A_{1n} \\ A_{21} & A_{22} & \cdots & A_{2n} \\ \vdots & \vdots & \cdots & \vdots \\ A_{m1} & A_{m2} & \cdots & A_{mn} \end{pmatrix}$$

- The transpose of A is

$$A^T = \begin{pmatrix} A_{11} & A_{21} & \cdots & A_{m1} \\ A_{12} & A_{22} & \cdots & A_{m2} \\ \vdots & \vdots & \cdots & \vdots \\ A_{1n} & A_{2n} & \cdots & A_{mn} \end{pmatrix}$$

- Matrix A is symmetric if
- Matrix norm: $\|A\|_p = \max \|Ax\|_p$ for $\|x\|_p = 1$
- $p=1 \rightarrow$ L1 norm
- $p=2 \rightarrow$ Frobenius norm
- $p=\text{infinity} \rightarrow$ L-infinity norm.

Matrix norms

Matrix norms are functions $f : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}$ that satisfy the same properties as vector norms. Let $A \in \mathbb{R}^{m \times n}$. Here are a few examples of matrix norms:

- The Frobenius norm: $\|A\|_F = \sqrt{\text{Tr}(A^T A)} = \sqrt{\sum_{i,j} A_{i,j}^2}$
- The sum-absolute-value norm: $\|A\|_{sav} = \sum_{i,j} |A_{i,j}|$
- The max-absolute-value norm: $\|A\|_{max} = \max_{i,j} |A_{i,j}|$

Basic concepts

We will use lower case letters for vectors The elements are referred by x_i .

- Vector dot (inner) product:

$$x^T y \in \mathbb{R} = \begin{bmatrix} x_1 & x_2 & \cdots & x_n \end{bmatrix} \begin{bmatrix} y_1 \\ x_2 \\ \vdots \\ y_n \end{bmatrix} = \sum_{i=1}^n x_i y_i.$$

If $u \cdot v = 0$, $\|u\|_2 \neq 0$, $\|v\|_2 \neq 0 \rightarrow u$ and v are *orthogonal*

If $u \cdot v = 0$, $\|u\|_2 = 1$, $\|v\|_2 = 1 \rightarrow u$ and v are *orthonormal*

- Vector outer product (direct product):

$$xy^T \in \mathbb{R}^{m \times n} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} \begin{bmatrix} y_1 & y_2 & \cdots & y_n \end{bmatrix} = \begin{bmatrix} x_1 y_1 & x_1 y_2 & \cdots & x_1 y_n \\ x_2 y_1 & x_2 y_2 & \cdots & x_2 y_n \\ \vdots & \vdots & \ddots & \vdots \\ x_m y_1 & x_m y_2 & \cdots & x_m y_n \end{bmatrix}$$

Basic concepts

We will use upper case letters for matrices. The elements are referred by A_{ij} .

- **Matrix product:**

$$A \in \mathbb{R}^{m \times n} \quad B \in \mathbb{R}^{n \times p}$$

$$C = AB \in \mathbb{R}^{m \times p}$$

$$C_{ij} = \sum_{k=1}^n A_{ik} B_{kj}$$

e.g.

$$A = \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix}, B = \begin{pmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{pmatrix}$$
$$AB = \begin{pmatrix} a_{11}b_{11} + a_{12}b_{21} & a_{11}b_{12} + a_{12}b_{22} \\ a_{21}b_{11} + a_{22}b_{21} & a_{21}b_{12} + a_{22}b_{22} \end{pmatrix}$$

Special matrices

$$\begin{pmatrix} a & 0 & 0 \\ 0 & b & 0 \\ 0 & 0 & c \end{pmatrix} \text{ diagonal} \quad \begin{pmatrix} a & b & c \\ 0 & d & e \\ 0 & 0 & f \end{pmatrix} \text{ upper-triangular}$$

$$\begin{pmatrix} a & b & 0 & 0 \\ c & d & e & 0 \\ 0 & f & g & h \\ 0 & 0 & i & j \end{pmatrix} \text{ tri-diagonal} \quad \begin{pmatrix} a & 0 & 0 \\ b & c & 0 \\ d & e & f \end{pmatrix} \text{ lower-triangular}$$

$$\begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \text{ I (identity matrix)}$$

Basic concepts

- Transpose

e.g. $\begin{pmatrix} a \\ b \end{pmatrix}^T = (a \quad b)$

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix}^T = \begin{pmatrix} a & c \\ b & d \end{pmatrix}$$

- $(A^T)^T = A$
- $(AB)^T = B^T A^T$
- $(A + B)^T = A^T + B^T$

Linear independence

- A set of vectors is **linearly independent** if none of them can be written as a linear combination of the others.
- Vectors v_1, \dots, v_k are linearly independent if $c_1 v_1 + \dots + c_k v_k = 0$ implies $c_1 = \dots = c_k = 0$

$$\begin{pmatrix} | & | & | \\ v_1 & v_2 & v_3 \\ | & | & | \end{pmatrix} \begin{pmatrix} c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}$$

e.g. $\begin{pmatrix} 1 & 0 \\ 2 & 3 \\ 1 & 3 \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}$ $(u,v)=(0,0)$, i.e. the columns are linearly independent.

$$x_1 = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \quad x_2 = \begin{bmatrix} 4 \\ 1 \\ 5 \end{bmatrix} \quad x_3 = \begin{bmatrix} 2 \\ -3 \\ -1 \end{bmatrix} \quad x_3 = -2x_1 + x_2$$

Span of a vector space

- If all vectors in a vector space may be expressed as linear combinations of a set of vectors v_1, \dots, v_k , then v_1, \dots, v_k **spans** the space.
- The cardinality of this set is the **dimension** of the vector space.

e.g.

$$\begin{pmatrix} 2 \\ 2 \\ 2 \end{pmatrix} = 2 \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} + 2 \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} + 2 \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

- A **basis** is a maximal set of linearly independent vectors and a minimal set of spanning vectors of a vector space

Rank of a Matrix

- $\text{rank}(A)$ (the rank of a m -by- n matrix A) is

The maximal number of linearly independent columns

=The maximal number of linearly independent rows

=The dimension of $\text{col}(A)$

=The dimension of $\text{row}(A)$

- If A is n by m , then

- $\text{rank}(A) \leq \min(m, n)$

- If $n = \text{rank}(A)$, then A has full row rank

- If $m = \text{rank}(A)$, then A has full column rank

$$\begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \quad \begin{pmatrix} 2 & 1 \\ 4 & 2 \end{pmatrix}$$

Inverse of a matrix

- Inverse of a square matrix A , denoted by A^{-1} is the *unique* matrix s.t.
 - $AA^{-1}=A^{-1}A=I$ (identity matrix)
- If A^{-1} and B^{-1} exist, then
 - $(AB)^{-1} = B^{-1}A^{-1}$,
 - $(A^T)^{-1} = (A^{-1})^T$
- For orthonormal matrices
- For diagonal matrices

$$\mathbf{A}^{-1} = \mathbf{A}^T$$

$$\mathbf{D}^{-1} = \text{diag}\{d_1^{-1}, \dots, d_n^{-1}\}$$

Matrix-vector multiplication

$$S = \begin{bmatrix} 3 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

has eigenvalues 3, 2, 0 with
corresponding eigenvectors

$$v_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}$$

$$v_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}$$

$$v_3 = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

On each eigenvector, S acts as a multiple of the identity matrix: but as a different multiple on each.

Any vector (say $x = \begin{pmatrix} 2 \\ 4 \\ 6 \end{pmatrix}$) can be viewed as a combination of the eigenvectors: $x = 2v_1 + 4v_2 + 6v_3$

Matrix-vector multiplication

- Thus a matrix-vector multiplication such as Sx (S , x as in the previous slide) can be rewritten in terms of the eigenvalues/vectors:

$$Sx = S(2v_1 + 4v_2 + 6v_3)$$

$$Sx = 2Sv_1 + 4Sv_2 + 6Sv_3 = 2\lambda_1v_1 + 4\lambda_2v_2 + 6\lambda_3v_3$$

- Even though x is an arbitrary vector, the action of S on x is determined by the eigenvalues/vectors.
- Suggestion: the effect of “small” eigenvalues is small.

Matrix calculus

	Scalar	Vector	Matrix
Scalar	$\frac{dy}{dx}$	$\frac{d\mathbf{y}}{dx} = \left[\frac{\partial y_i}{\partial x} \right]$	$\frac{d\mathbf{Y}}{dx} = \left[\frac{\partial y_{ij}}{\partial x} \right]$
Vector	$\frac{dy}{d\mathbf{x}} = \left[\frac{\partial y}{\partial x_j} \right]$	$\frac{d\mathbf{y}}{d\mathbf{x}} = \left[\frac{\partial y_i}{\partial x_j} \right]$	
Matrix	$\frac{dy}{d\mathbf{X}} = \left[\frac{\partial y}{\partial x_{ji}} \right]$		

By Thomas Minka. Old and New Matrix Algebra Useful for Statistics

Examples

$$\frac{\partial \mathbf{x}^T \mathbf{a}}{\partial \mathbf{x}} = \frac{\partial \mathbf{a}^T \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}$$

$$\frac{\partial \mathbf{a}^T \mathbf{X} \mathbf{b}}{\partial \mathbf{X}} = \mathbf{a} \mathbf{b}^T$$

$$\frac{\partial \mathbf{a}^T \mathbf{X}^T \mathbf{b}}{\partial \mathbf{X}} = \mathbf{b} \mathbf{a}^T$$

$$\frac{\partial \mathbf{a}^T \mathbf{X} \mathbf{a}}{\partial \mathbf{X}} = \frac{\partial \mathbf{a}^T \mathbf{X}^T \mathbf{a}}{\partial \mathbf{X}} = \mathbf{a} \mathbf{a}^T$$

$$\frac{\partial \mathbf{x}^T \mathbf{B} \mathbf{x}}{\partial \mathbf{x}} = (\mathbf{B} + \mathbf{B}^T) \mathbf{x}$$

<http://matrixcookbook.com/>

Eigenvalues & Eigenvectors

- A scalar λ is an eigenvalue of an $n \times n$ matrix A if there is a nonzero vector x such that $Ax = \lambda x$.
 - Vector x is called an eigenvector.
- Matrix A is symmetric positive definite (SPD) if $A^T = A$ and all its eigenvalues are positive.
- If A has n linearly independent eigenvectors, A can have the eigen-decomposition: $A = X\Lambda X^{-1}$.
 - Λ is diagonal with eigenvalues as its diagonal elements
 - Column vectors of X are corresponding eigenvectors

Eigenvalues & Eigenvectors

- **Eigenvectors** (for a square $m \times m$ matrix \mathbf{S})

$$\mathbf{S}\mathbf{v} = \lambda\mathbf{v}$$

(right) eigenvector

eigenvalue

$$\mathbf{v} \in \mathbb{R}^m \neq \mathbf{0}$$

$$\lambda \in \mathbb{R}$$

- How many eigenvalues are there at most?

Example

$$\begin{pmatrix} 6 & -2 \\ 4 & 0 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \end{pmatrix} = \begin{pmatrix} 2 \\ 4 \end{pmatrix} = 2 \begin{pmatrix} 1 \\ 2 \end{pmatrix}$$

$$\mathbf{S}\mathbf{v} = \lambda\mathbf{v} \iff (\mathbf{S} - \lambda\mathbf{I})\mathbf{v} = \mathbf{0}$$

only has a non-zero solution if $|\mathbf{S} - \lambda\mathbf{I}| = 0$

this is a m -th order equation in λ which can have **at most m distinct solutions** (roots of the characteristic polynomial) - can be complex even though \mathbf{S} is real.

Eigenvalues & Eigenvectors

For symmetric matrices, eigenvectors for distinct eigenvalues are **orthogonal**

$$Sv_{\{1,2\}} = \lambda_{\{1,2\}}v_{\{1,2\}}, \text{ and } \lambda_1 \neq \lambda_2 \Rightarrow v_1 \bullet v_2 = 0$$

All eigenvalues of a real symmetric matrix are **real**.

for complex λ , if $|S - \lambda I| = 0$ and $S = S^T \Rightarrow \lambda \in \mathfrak{R}$

All eigenvalues of a **positive semidefinite** matrix are **non-negative**

$$\forall w \in \mathfrak{R}^n, w^T S w \geq 0, \text{ then if } S v = \lambda v \Rightarrow \lambda \geq 0$$

Example

- Let $S = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$ ← Real, symmetric.

- Then $S - \lambda I = \begin{bmatrix} 2 - \lambda & 1 \\ 1 & 2 - \lambda \end{bmatrix} \Rightarrow (2 - \lambda)^2 - 1 = 0.$

- The eigenvalues are 1 and 3 (nonnegative, real).

- The eigenvectors are orthogonal (and real):

$$\begin{pmatrix} 1 \\ -1 \end{pmatrix} \quad \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$

Plug in these values and solve for eigenvectors.

Eigen/diagonal Decomposition

- Let $S \in \mathbb{R}^{m \times m}$ be a **square** matrix with **m linearly independent eigenvectors** (a “non-defective” matrix)

- **Theorem:** Exists an **eigen decomposition**

- (cf. matrix diagonalization theorem) $S = U \Lambda U^{-1}$ *diagonal*

Unique
for
distinct
eigen-
values

- Columns of U are **eigenvectors** of S
- Diagonal elements of Λ are **eigenvalues** of S

$$\Lambda = \text{diag}(\lambda_1, \dots, \lambda_m), \quad \lambda_i \geq \lambda_{i+1}$$

Diagonal decomposition: why/how

Let **U** have the eigenvectors as columns: $U = \begin{bmatrix} v_1 & \dots & v_n \end{bmatrix}$

Then, **SU** can be written

$$SU = S \begin{bmatrix} v_1 & \dots & v_n \end{bmatrix} = \begin{bmatrix} \lambda_1 v_1 & \dots & \lambda_n v_n \end{bmatrix} = \begin{bmatrix} v_1 & \dots & v_n \end{bmatrix} \begin{bmatrix} \lambda_1 & & \\ & \dots & \\ & & \lambda_n \end{bmatrix}$$

Thus **$SU=U\Lambda$** , or **$U^{-1}SU=\Lambda$**

And **$S=U\Lambda U^{-1}$** .

Diagonal decomposition - example

Recall $S = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}; \lambda_1 = 1, \lambda_2 = 3.$

The eigenvectors $\begin{pmatrix} 1 \\ -1 \end{pmatrix}$ and $\begin{pmatrix} 1 \\ 1 \end{pmatrix}$ form $U = \begin{bmatrix} 1 & 1 \\ -1 & 1 \end{bmatrix}$

Inverting, we have $U^{-1} = \begin{bmatrix} 1/2 & -1/2 \\ 1/2 & 1/2 \end{bmatrix}$

Recall
 $UU^{-1} = I.$

Then, $S = U\Lambda U^{-1} = \begin{bmatrix} 1 & 1 \\ -1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} 1/2 & -1/2 \\ 1/2 & 1/2 \end{bmatrix}$

Example continued

Let's divide \mathbf{U} (and multiply \mathbf{U}^{-1}) by $\sqrt{2}$

$$\text{Then, } \mathbf{S} = \underbrace{\begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ -1/\sqrt{2} & 1/\sqrt{2} \end{bmatrix}}_{\mathbf{U}} \underbrace{\begin{bmatrix} 1 & 0 \\ 0 & 3 \end{bmatrix}}_{\mathbf{\Lambda}} \underbrace{\begin{bmatrix} 1/\sqrt{2} & -1/\sqrt{2} \\ 1/\sqrt{2} & 1/\sqrt{2} \end{bmatrix}}_{\mathbf{U}^T}$$

Symmetric Eigen Decomposition

- If $S \in \mathbb{R}^{m \times m}$ is a **symmetric** matrix:
- **Theorem**: Exists a (unique) **eigen decomposition** $S = Q\Lambda Q^T$
- where Q is **orthogonal**:
 - $Q^{-1} = Q^T$
 - Columns of Q are normalized eigenvectors
 - Columns are orthogonal.
 - (everything is real)

Spectral decomposition

- If A is real and symmetric, all its eigenvalues are real, and there are n orthogonal eigenvectors.
- The spectral decomposition of a symmetric matrix A is $A=Q\Lambda Q^T$.
 - Λ is diagonal with eigenvalues as its diagonal elements
 - Q is orthogonal, i.e. $Q^T Q = Q Q^T = I$.
 - Column vectors of Q are corresponding eigenvectors.

Exercise

- Examine the symmetric eigen decomposition, if any, for each of the following matrices:

$$\begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 \\ -2 & 3 \end{bmatrix}$$

$$\begin{bmatrix} 2 & 2 \\ 2 & 4 \end{bmatrix}$$

Singular Value Decomposition (SVD)

- The singular values of an $m \times n$ A are the square roots of the eigenvalues of $A^T A$.
- Any matrix A can have the singular value decomposition (SVD): $A = U \Sigma V^T$.
 - Σ is diagonal with singular values as its elements.
 - U and V are orthogonal matrices.
 - The column vectors of U are called left singular vectors of A ; the column vectors of V is called the right singular vector of A .

Singular Value Decomposition

For an $m \times n$ matrix \mathbf{A} of rank r there exists a factorization (Singular Value Decomposition = **SVD**) as follows:

$$\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$$

$m \times m$ $m \times n$ V is $n \times n$

The columns of \mathbf{U} are orthogonal eigenvectors of $\mathbf{A}\mathbf{A}^T$.

The columns of \mathbf{V} are orthogonal eigenvectors of $\mathbf{A}^T\mathbf{A}$.

Eigenvalues $\lambda_1 \dots \lambda_r$ of $\mathbf{A}\mathbf{A}^T$ are the eigenvalues of $\mathbf{A}^T\mathbf{A}$.

$$\sigma_i = \sqrt{\lambda_i}$$
$$\mathbf{\Sigma} = \text{diag}(\sigma_1 \dots \sigma_r)$$

← Singular values.

Singular Value Decomposition

- Illustration of SVD dimensions and sparseness

$$\underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_A = \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & & & \\ & \bullet & & & \\ & & \bullet & & \\ & & & \bullet & \\ & & & & \bullet \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_{V^T}$$

$$\underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_A = \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & & & \\ & \bullet & & & \\ & & \bullet & & \\ & & & \bullet & \\ & & & & \bullet \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{V^T}$$

SVD example

Let $A = \begin{bmatrix} 1 & -1 \\ 0 & 1 \\ 1 & 0 \end{bmatrix}$

Thus $m=3, n=2$. Its SVD is

$$\begin{bmatrix} 0 & 2/\sqrt{6} & 1/\sqrt{3} \\ 1/\sqrt{2} & -1/\sqrt{6} & 1/\sqrt{3} \\ 1/\sqrt{2} & 1/\sqrt{6} & -1/\sqrt{3} \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & \sqrt{3} \\ 0 & 0 \end{bmatrix} \begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix}$$

Typically, the singular values arranged in decreasing order.

Low-rank Approximation

- SVD can be used to compute optimal **low-rank approximations**.
- Approximation problem: Find \mathbf{A}_k of rank k such that

$$\mathbf{A}_k = \min_{X: \text{rank}(X)=k} \|\mathbf{A} - \mathbf{X}\|_F \longleftarrow \text{Frobenius norm}$$

$$\|\mathbf{A}\|_F \equiv \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}.$$

\mathbf{A}_k and \mathbf{X} are both $m \times n$ matrices.

Typically, want $k \ll r$.

Low-rank Approximation

- Solution via SVD

$$A_k = U \operatorname{diag}(\sigma_1, \dots, \sigma_k, \underbrace{0, \dots, 0}_{\substack{\text{set smallest } r-k \\ \text{singular values to zero}}}) V^T$$

$$\underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{A_k} = \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & \\ & \bullet & \\ & & \bullet \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{V^T}$$

$$A_k = \sum_{i=1}^k \sigma_i u_i v_i^T \longleftarrow \text{column notation: sum of rank 1 matrices}$$

Approximation error

- How good (bad) is this approximation?
- It's the best possible, measured by the Frobenius norm of the error:

$$\min_{X: \text{rank}(X)=k} \|A - X\|_F = \|A - A_k\|_F = \sigma_{k+1}$$

where the σ_i are ordered such that $\sigma_i \geq \sigma_{i+1}$.

Suggests why Frobenius error drops as k increased.

C. Eckart, G. Young, *The approximation of a matrix by another of lower rank.*
Psychometrika, 1, 211-218, 1936.

SVD Low-rank approximation


- Given a term-doc matrix A of $m=50000$, $n=10$ million (and rank close to 50000).
- We can construct an approximation A_{100} with rank 100.
 - Of all rank 100 matrices, it would have the lowest Frobenius error.
- Great ... but why would we??
- **Answer:** *Latent Semantic Document Indexing, Movie Recommendation,...*

Eigen Value Decomposition

- Any **symmetric** matrix A can be decomposed as $A=UDU^T$, where
 - D is diagonal, with $d=\text{rank}(A)$ non-zero elements
 - The first d rows of U are orthogonal basis for $\text{col}(A)=\text{row}(A)$

■ Re-interpreting Ab

- First stretch b along the direction of u_1 by d_1 times
- Then further stretch it along the direction of u_2 by d_2 times



Summary of SVD

- Any general matrix A can be decomposed as $A=UDV^T$, where
 - D is a diagonal matrix, with $d=\text{rank}(A)$ non-zero elements
 - The first d rows of U are orthogonal basis for $\text{col}(A)$
 - The first d rows of V are orthogonal basis for $\text{row}(A)$
- Applications of the SVD
 - Matrix Pseudoinverse
 - Low-rank matrix approximation

LU decomposition

- The LU decomposition with pivoting of matrix A is $PA=LU$
 - P is a permutation matrix
 - L is lower triangular; U is upper triangular.
- The linear system $Ax=b$ can be solved by
 1. Perform LU decompose $PA=LU$
 2. Solve $Ly=Pb$
 3. Solve $Ux=y$

Cholesky decomposition

- For an SPD matrix A , there exists the Cholesky decomposition $P^TAP = LL^T$
 - P is a permutation matrix
 - L is a lower triangular matrix
- If A is not SPD, the LBL decomposition can be used: $P^TAP = LBL^T$
 - B is a block diagonal matrix with blocks of dimension 1 or 2.

Subspaces, QR decomposition

- The null space of an $m \times n$ matrix A is

$$\mathbf{Null}(A) = \{w \mid Aw = 0, w \neq 0\}$$

- The range of A is $\mathbf{Range}(A) = \{w \mid w = Av, \forall v\}$.
- Fundamental of linear algebra:

$$\mathbf{Null}(A) \oplus \mathbf{Range}(A^T) = \mathbb{R}^n$$

- Matrix A has the QR decomposition $AP = QR$
 - P is permutation matrix; Q is an orthogonal matrix;
 R is an upper triangular matrix.

Recap: linear algebra

- Vectors and matrix
- Eigenvalue and eigenvector
- Singular value decomposition
- LU decomposition and Cholesky decomposition
- Subspaces and QR decomposition

End of Math Review-1

- Read the 3 Stanford University (CS229) lecture notes provide on Wattle.
- Read an old (2019) ENGN8535 lecture notes on Wattle.

Tomorrow's tutorial (zoom based, at lecture time)

- Python Tutorial
- Numpy
- Github and GitLab (for CLab submission).
- Please start to practice Python programming. You will need it for CLab-1.

End of today's lecture.